

🛡 PROJECT: NIST VULNERABILITY EXTRACTOR

NIST Vulnerability Information Extractor

A Transformer-based Named Entity Recognition (NER) system for automated cybersecurity intelligence extraction using DeBERTa architecture.

📅 December 2025 🗂️ CyNER2.0 Augmented 🌐 DeBERTa v3 Base

Project Team

👤 Ishaan Singh
LCI2022025

👤 Aun Abbas
LCI2022014

👤 Shashank Agarwal
LCI2022045

👤 Aaditya Mishra
LCI2022001

👤 Disha Yadav
LCI2022021

Introduction: Automating CVE Extraction



The Challenge: Unstructured Data

Cybersecurity ecosystems rely on CVE entries, but descriptions are unstructured text. Manual analysis is slow, error-prone, and cannot scale with the volume of modern threats.



The Goal: Automated Extraction

Develop a system to automatically identify and classify critical entities: [Products](#), [OS](#), [Attack Vectors](#), and [Impact Effects](#).



The Approach: Transformer NER

We built a Named Entity Recognition (NER) system using **DeBERTa v3**, fine-tuned on an augmented cybersecurity dataset (CyNER2.0) to handle technical jargon.



The Outcome: Scalable Intelligence

A robust pipeline that transforms raw text into structured JSON, enabling immediate integration with SIEM tools, SOC dashboards, and vulnerability scanners.

WORKFLOW VISUALIZATION



INPUT: RAW TEXT

"Buffer overflow in OpenSSL 1.0.2..."



DeBERTa NER Model

Token Classification

OUTPUT: STRUCTURED JSON

Product: OpenSSL

Version: 1.0.2

Vuln: Buffer Overflow

Objectives & Use Cases

◎ Core Objectives

- ▶ Achieve **high-accuracy Named Entity Recognition (NER)** for specialized cybersecurity entities within unstructured text.
- ▶ Develop robustness to handle **rare technical terms** and product names often missed by general-purpose models.
- ▶ Ensure scalability to process **massive volumes** of CVE descriptions and security advisories in real-time.

📦 Key Deliverables

 Fine-tuned DeBERTa Model

 End-to-End Inference Pipeline

 Streamlit & Gradio Apps

 Comprehensive Technical Report

💻 Application Ecosystem



SIEM Enrichment

Automatically tag incoming logs with affected products and vulnerabilities to prioritize alerts in Security Information and Event Management systems.



Vulnerability Scanners

Parse scan reports to extract standardized entity data for cross-referencing against internal asset inventories.



SOC Dashboards

Visualize trends in attack vectors and affected operating systems extracted from daily threat feeds.



Threat Intelligence

Build knowledge graphs linking threat actors to specific product versions and vulnerability types from unstructured reports.

Dataset: CyNER2.0 Augmented

 DATA SOURCE

Dataset Source

[PranavaKailash/CyNER2.0_augmented_dataset](#)

✓ Hugging Face ★ High Quality 💫 NER Task

An improved, enlarged version of the original CyNER 2.0 dataset specifically designed for training robust cybersecurity entity recognition models.

Dataset Composition



CVE Descriptions

Sourced directly from MITRE CVE and NIST NVD databases covering wide range of vulnerabilities.



Vendor Advisories

Security bulletins from major software vendors (Microsoft, Cisco, RedHat) providing diverse technical context.



Manual Annotations

Expert-verified labels ensuring high ground-truth quality for training and evaluation.

Synthetic Augmentations



Original CVEs
Vendor Advisories
Synthetic Data
Manual Annotation

AUGMENTATION TECHNIQUES

Synonym Replacement

"Execute code" → "Run arbitrary commands"

VOCABULARY

Version Swapping

"Apache 2.4.49" → "Apache 2.4.50-rc1"

ROBUSTNESS

Template Generation

Synthesizing new sentence structures for unseen formats

STRUCTURE

Product-OS Combos

Randomizing "PostgreSQL on Windows" vs "PostgreSQL on Fedora"

DIVERSITY

Entity Labels: BIO Schema Classes

8 TOTAL CLASSES


ENTITY LABEL	DESCRIPTION & MEANING	REAL-WORLD EXAMPLES
PRODUCT	Software, hardware, or firmware names affected by the vulnerability.	OpenSSL, Apache HTTPD, Cisco ASA, Zoom Client
OS	Operating systems mentioned as the platform or environment.	Ubuntu 18.04, Windows Server 2016, macOS Big Sur
ENVIRONMENT	Runtime environments, containers, or virtualization contexts.	Docker Container, Kubernetes pod, VMware vSphere
VERSION / PATCH	Specific version numbers, builds, or patch identifiers.	prior to 1.0.2g, 7.5.2, build 2021-04, KB 5001330
ATTACK_VECTOR	The method or pathway used to deliver the exploit.	Remote attacker, via network, physically present user
PREREQUISITE	Conditions required before exploitation is possible.	Requires login, authenticated user, with admin privileges
EFFECT	The consequence or impact of successful exploitation.	Code execution, privilege escalation, denial of service
CWE/VULN_TYPE	The category or technical nature of the weakness.	Buffer overflow, SQL injection, Cross-site scripting (XSS)

Annotation Format

Data follows the **CoNLL** token-tagging format using BIO (Beginning, Inside, Outside) notation.

OpenSSL	B-PRODUCT
1.0.2	B-VERSION
on	O
Linux	B-OS

LABEL DISTRIBUTION



Dataset: Format & Augmentation



CoNLL Token-Tagging Format

Standard BIO (Begin-Inside-Outside) schema used for token-level annotation. Each token is aligned with a specific label, ensuring granular entity boundaries even for complex technical terms.

OpenSSL B-PRODUCT

1.0.0g B-VERSION

buffer B-VULN

overflow I-VULN



Token-Level Annotation Alignment

Annotations are strictly aligned to subword tokens generated by the tokenizer (WordPiece/BPE). This prevents misalignment issues common in cybersecurity texts containing special characters (e.g., CVE-2023-1234, mod_ssl).



Purpose of Augmentation

Cybersecurity vocabulary is niche and rapidly evolving. Augmentation addresses:

- Handling rare product names (e.g., legacy or custom software)
- Improving recognition of long entity spans (e.g., "arbitrary code execution via buffer overflow")
- Increasing robustness for unseen CVE formats and varied phrasing



Robustness to Phrasing

By training on augmented sentences (synonyms, version swaps, templates), the model generalizes better to real-world reports that differ from standard NVD descriptions.



DATA_FORMAT_VISUALIZATION

Token	BIO Label
Apache	B-PRODUCT
Tomcat	I-PRODUCT
9	B-VERSION
.	I-VERSION
0	I-VERSION
allows	0
RCE	B-EFFECT



AUGMENTATION_STRATEGY

ORIGINAL "Ubuntu 20.04"	→	AUGMENTED "Debian 11"
ORIGINAL "Code Execution"	→	AUGMENTED "Arbitrary command run"

Model Architecture

DEBERTA V3 BASE

Core Specifications


Base Model	DeBERTa-v3-base
Fine-Tuned Ver.	cyner_ner_model
Task Type	Token Classification
Output Head	Linear + Softmax

Why DeBERTa?

- ▶ **Disentangled Attention:** Unlike BERT, it separates content and position embeddings, allowing better modeling of relative positions crucial for technical syntax.
- ▶ **Enhanced Mask Decoder:** Incorporates absolute positions in the decoding layer to improve context awareness for masked token prediction.
- ▶ **Domain Adaptation:** Superior handling of "context-dependent" cybersecurity terms (e.g., distinguishing "remote" as vector vs. location).

END-TO-END PIPELINE VISUALIZATION

INPUT PROCESSING
Token Embeddings + Position Embeddings



CONTEXT REFINEMENT
Enhanced Mask Decoder (EMD)



OUTPUT HEAD
Linear Layer (Hidden → Num_Labels)
Softmax Activation

B-PROD I-PROD O B-VULN

DeBERTa Innovations for Technical Language



Disentangled Attention Mechanism

Unlike BERT which sums embeddings, DeBERTa separates **content embeddings** from **position embeddings**. Two distinct attention scores are computed: content-to-content and content-to-position.

Content Vector

Position Vector



Explicit Relative Positioning

The model understands the *relative distance* between tokens rather than just absolute positions. This is critical for parsing complex technical syntax where word order defines relationships (e.g., "Apache 2.4" vs "2.4 Apache").



Enhanced Mask Decoder (EMD)

Incorporates absolute position information right before the softmax layer during pre-training. This significantly improves the model's ability to predict masked tokens based on both context and exact location.



Domain Context Preservation

These architectural choices help preserve the specific meaning of polysemous words in cybersecurity. Terms like "remote" (attack vector) vs. "remote" (adjective) are distinguished with higher accuracy.

ATTENTION_COMPARISON



Traditional BERT

Content + Position

Single Vector

DeBERTa Approach

CONTENT STREAM

POSITION STREAM

DISENTANGLING ATTENTION MATRICES

Content-to-Content

Content-to-Position

RESULT:

Superior handling of specialized syntax found in CVE descriptions and technical advisories.

Tokenization & Embedding Pipeline

>_ Tokenizer Mechanics

The tokenizer transforms raw text into numeric representations essential for the model. It produces three key outputs:

- `token_ids` Numerical ID for each token in vocab
- `attn_mask` Binary mask (1=content, 0=padding)
- `type_ids` Segment ID (usually 0 for single seq)

<# Subword Handling (WordPiece/BPE)

Cybersecurity texts contain rare technical terms. Subword tokenization breaks these down into meaningful components.

EXAMPLE: "OPENSSL VULNERABILITY"

Open ##SSL vulnerability

Token 1 Token 2 Token 3

Label Alignment: BIO tags are expanded to cover all sub-tokens (e.g., B-PRODUCT, I-PRODUCT).

≡ Embedding Architecture

INPUT TOKEN

"##SSL"

A Token Embedding

📍 Position Embedding

▣ Segment Embedding



Final Input Representation

Vector sum passed to Encoder Layers



Encoder & Classification Head

Transformer Encoder Layers

The core processing unit consists of stacked transformer layers designed to capture deep contextual relationships within technical cybersecurity text.

LAYER COMPONENTS

- ✓ Multi-head Self-Attention
- ✓ Feed-Forward Network
- ✓ Residual Connections
- ✓ Layer Normalization

- ▶ **Contextual Mapping:** Self-attention mechanisms enable the model to relate distant tokens, such as linking a product name "OpenSSL" at the start of a sentence to a version number "1.0.2g" at the end.
- ▶ **Domain Adaptation:** Crucial for disambiguating terms like "exploit" (noun vs. verb) or recognizing "buffer overflow" as a specific vulnerability type (CWE).

Token Classification Head

The final stage projects the high-dimensional hidden states into specific entity class probabilities for every token in the sequence.

Linear Projection

```
Logits = Linear(H_token)
Projects 768-dim hidden state → num_labels dimension.
```

Softmax Activation

```
P(class | token) = Softmax(Logits)
Outputs probability distribution across all BIO tags (B-OS, I-PRODUCT, O, etc.).
```

Token Classification Logic Flow

Transformer Output

 $H \in \mathbb{R}^{d}$

Contextualized Embedding (768-dim)

Linear Layer

Weight Matrix (W) $\times H + \text{Bias } (b)$

Dimensionality Reduction

Raw Logits [N_Classes]

Softmax Function


B-OS

B-PRODUCT

O

ArgMax Selection



Training Pipeline Overview

[END-TO-END FLOW](#)

Training Hyperparameters

[MODEL CONFIGURATION](#)

PARAMETER	VALUE / SETTING	DESCRIPTION & CONTEXT
Base Model	DeBERTa-v3-base	Pre-trained transformer foundation optimized for NLU tasks.
Max Length	256 tokens	Padding/truncation limit covering most CVE descriptions.
Batch Size	8 – 16	Dependent on GPU VRAM availability (typically 16GB+).
Optimizer	AdamW	Adam with weight decay fix; standard for transformers.
Learning Rate	2e-5	Low rate to prevent catastrophic forgetting during fine-tuning.
Epochs	3 – 6	Early stopping applied if validation loss plateaus.
Scheduler	Linear Warmup	Gradual LR increase for first 10% steps, then linear decay.
Loss Function	Cross Entropy	Token-level loss calculation with padding masks applied.



Class Imbalance & Metrics



The Imbalance Challenge

Cybersecurity entities are naturally skewed. Common entities like OS (e.g., Windows) appear frequently, while critical but specific entities like EFFECT (e.g., privilege escalation) are much rarer.

HIGH: PRODUCT

LOW: EFFECT



Mitigation Strategies

- ✓ **Weighted Loss:** Assign higher penalty weights to scarcity classes during backpropagation.
- ✓ **Oversampling:** Artificially increase frequency of rare entities in training batches.
- ✓ **Augmentation:** Generate synthetic sentences specifically for underrepresented classes.



Evaluation Metrics

Standard accuracy is misleading for imbalanced data. We focus on entity-level performance:

PRECISION

Correctness of positive predictions

RECALL

Coverage of actual positive cases

F1-SCORE

Harmonic mean of Precision & Recall

SPAN ACCURACY

Exact match of full entity boundaries

DISTRIBUTION_VISUALIZATION



Entity Frequency: Before vs After Augmentation



WEIGHT ADJUSTMENT LOGIC

Loss(Product)	x 1.0
Loss(Effect)	x 2.5
Loss(Attack_Vec)	x 1.8

Inference Pipeline

⚡ REAL-TIME EXTRACTION


≡ RAW TEXT →

[TOKENS] →

[LOGITS] →

{ ENTITIES } →

</> JSON OUTPUT

Inference Example


〈/〉 LIVE TRANSFORMATION

1 Raw Input Text

"A buffer overflow in OpenSSL on Ubuntu 18.04 allows attackers to execute arbitrary code."



2 Token Classification (Model Output)



● PRODUCT

● OS

● VULN_TYPE

EFFECT

3 Structured JSON Output

JSON Format

```
[
  {
    "entity": "PRODUCT",
    "value": "OpenSSL",
    "confidence": 0.9942
  },
  {
    "entity": "OS",
    "value": "Ubuntu 18.04",
    "confidence": 0.9876
  },
  {
    "entity": "VULN_TYPE",
    "value": "buffer overflow",
    "confidence": 0.9631
  },
  {
    "entity": "EFFECT",
    "value": "execute arbitrary code",
    "confidence": 0.9205
  }
]
```

↗ Post-Processing Applied:
 Sub-tokens like "Open" + "##SSL" merged.
 "Ubuntu" + "18" + "." + "04" grouped into single OS entity.

Frontend Applications



Streamlit Implementation

LOCAL DEV

Designed for rapid local execution and development testing.

Provides a lightweight interface for direct model interaction.

- ✓ Simple text box input mechanism
- ✓ Clear entities table visualization
- ✓ Automatic entity grouping display
- ✓ Easy local deployment



Gradio Interface

PRODUCTION

Polished web interface suitable for public demonstrations and HuggingFace Spaces deployment. Features enhanced responsiveness.

- ✓ Live inference updates
- ✓ Fast prototyping capabilities
- ✓ Built-in API endpoint generation
- ✓ Seamless HuggingFace integration



UI_PREVIEW.exe

DEPLOYED



NIST Vulnerability Extractor

INPUT CVE DESCRIPTION

A buffer overflow vulnerability in OpenSSL 1.0.2 allows remote attackers to execute arbitrary code via...

EXTRACT ENTITIES

EXTRACTED ENTITIES

Entity Type	Value
PRODUCT	OpenSSL
VERSION	1.0.2
VULN_TYPE	Buffer Overflow
EFFECT	Execute Code

SYSTEM READY

v1.2.0-stable

Challenges & Solutions

ISSUE

RESOLUTION

Pipeline Import Failure

Runtime error when initializing the Hugging Face pipeline, causing
`ImportError: cannot import name 'pipeline'.`

```
ModuleNotFoundError: No module named 'transformers.pipelines'
```



Version Compatibility Fix

Identified version mismatch in the environment. Downgrading to a stable compatible version resolved the dependency conflict.



```
pip install transformers==4.53
```

Meta Tensor Error

Model weights failed to load correctly when initializing from cache, resulting in tensor shape mismatch or empty weights.

```
RuntimeError: Expected all tensors to be on the same device
```



Local Weight Loading

Forced loading from local directory to ensure complete weight initialization and correct device placement.



```
from_pretrained(local_dir, local_files_only=True)
```

Niche Product Detection

Low recall on obscure security tools (e.g., "OpenVAS", "Webmin") due to limited representation in training data.

```
Recall < 60%
```



Subword Augmentation

Implemented data augmentation with synthetic CVEs and leveraged DeBERTa's subword tokenization to generalize unseen terms.



```
F1 Score > 85%
```

Current System Limitations



Uncommon Vendor Products

The model occasionally misses niche or legacy vendor products not present in the augmented training set. While common vendors (Microsoft, Adobe, Cisco) are well-recognized, specialized industrial or obscure software may be overlooked.



Effect Wording Variance

Impact effects ("execute arbitrary code", "gain privilege") have high linguistic variance. Normalizing these extracted phrases into standard taxonomy remains challenging due to the diverse ways vulnerability impacts are described by different analysts.



Multi-Sentence Complexity

Performance degrades with long, multi-sentence CVE descriptions where context is separated by significant distance. Cross-sentence dependencies (referring to "the software" in the second sentence) can lead to entity resolution gaps.



Dataset Distribution

Although augmented, CyNER2.0 represents only a fraction of the total real-world CVE distribution. The model's generalization to completely novel attack vectors or emerging technologies (like new blockchain protocols) is limited by current training data.

PERFORMANCE_BOTTLENECKS



ERROR FREQUENCY DISTRIBUTION

Rare Products	12%
---------------	-----



Complex Syntax	8%
----------------	----



Context Loss	5%
--------------	----



CONSTRAINT IMPACT MATRIX

High

RECALL ON NICHE
SOFTWARE

Med

LONG-TEXT
COHERENCE

Low

STANDARD OS
DETECTION

None

INFERENCE SPEED

Current limitations primarily affect edge cases. Core functionality remains robust for >90% of standard NVD entries.

Future Roadmap & Conclusion

☒ PROJECT COMPLETE

📍 FUTURE ENHANCEMENTS



Relation Extraction

Link extracted entities to understand context (e.g., explicitly connecting "OpenSSL" to "Ubuntu 18.04").



CVSS Score Prediction

Add a classification head to predict vulnerability severity scores based on the description text.



Larger Corpora Integration

Expand training data with ExploitDB, Microsoft Security Advisories, and vendor disclosures.



Production Deployment

Containerize inference pipeline using Docker and expose via high-performance FastAPI endpoints.

☑ PROJECT SUMMARY

NIST Vulnerability Extractor

FULL-FLEDGED NLP SYSTEM

This project successfully demonstrates the application of modern Transformer-based NLP to the high-stakes domain of cybersecurity.



MODEL

DeBERTa-v3 Transformer Architecture



DATASET

CyNER2.0 Augmented Custom Dataset



INTERFACE

Streamlit & Gradio Frontends



READY FOR INTEGRATION

SIEM / SOC / THREAT INTEL

